Instructions -

Awesome—I've dropped a **ready-to-run starter project** into your workspace.

Quick start:

1. Download the README

2. Download requirements.txt

3. Download the DuckDB SQL pipeline

4. Download the Python runner

How to use it (5 steps):

Create a venv and install deps:

```
 cd otter_churn_project
python -m venv .venv && source .venv/bin/activate
pip install -r requirements.txt
```

1.
2. Put your three CSVs into `otter_churn_project/data_raw/` with these names:

    - `oa_churn_requests_sample.csv`

    - `oa_account_dimensions_sample.csv`

    - `oa_product_activity_sample.csv`

Run the end-to-end pipeline:
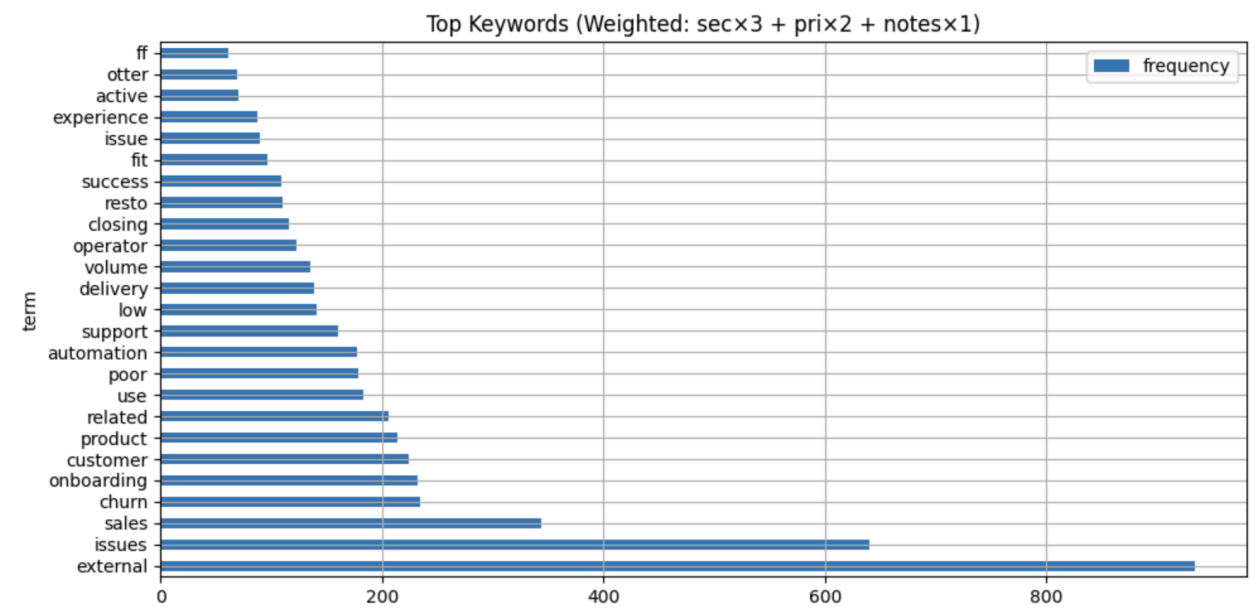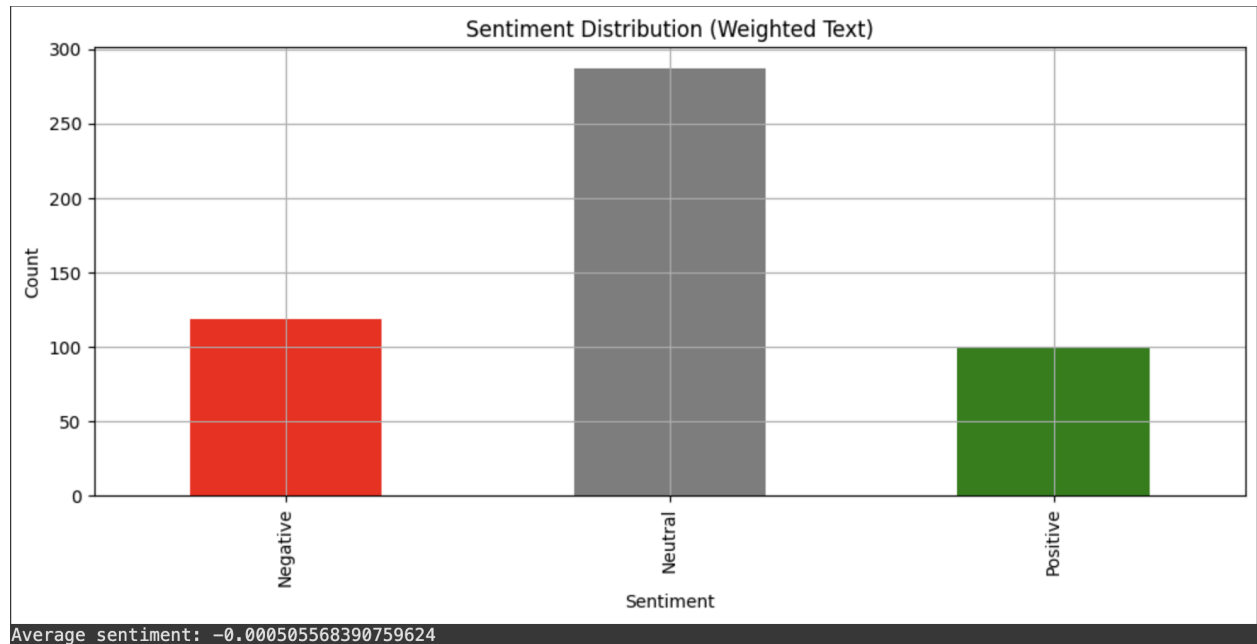
```
 python notebooks/otter_churn_starter.py
```

3. This executes the DuckDB SQL, builds features (7/30-day windows, adoption breadth, days since last use), and exports clean CSVs to `data_work/` for your dashboard.

4. Connect the exports to your dashboard:

- ○ Upload the CSVs in `data_work/` to Google Drive → open in Google Sheets.

- ○ In Looker Studio, add those Sheets as data sources and build the visuals (monthly churn trend, top reasons, reasons×segment, risk table).

5. Build your slides: screenshot key charts from Looker Studio or export PNGs from Python and drop them into Google Slides.
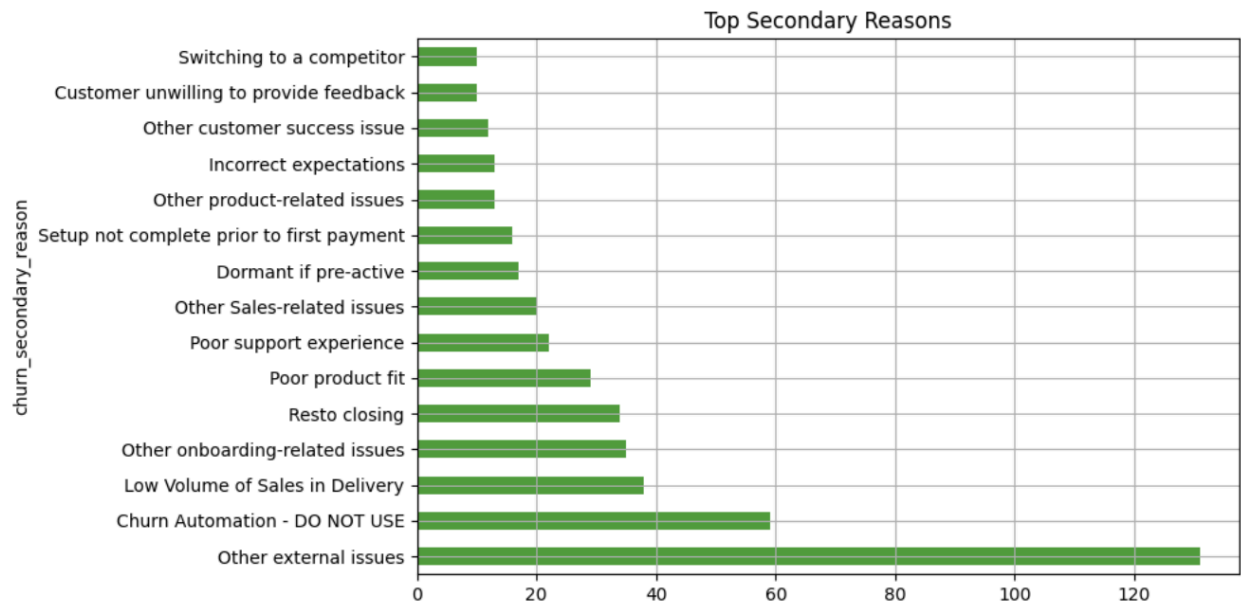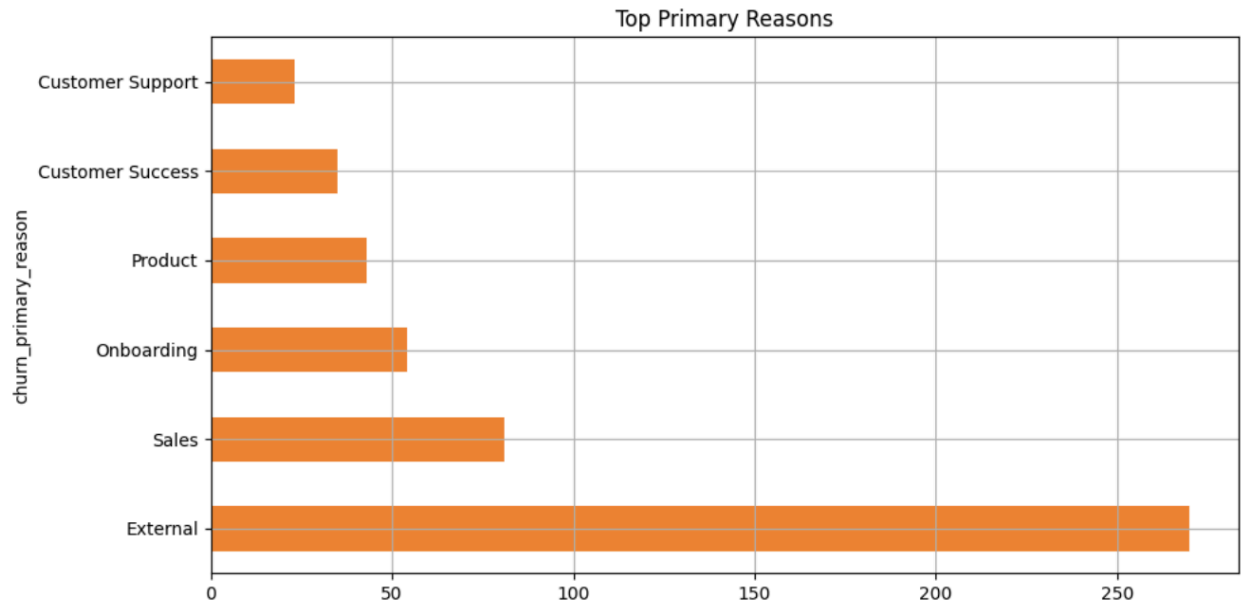
What's inside the files:

- ● `sql/otter_churn_duckdb.sql`

  - ○ Reads CSVs in place with `read_csv_auto`, profiles, builds:

    - ■ `dim_accounts`, `churn_events`, `fct_activity`

    - ■ `fct_prechurn` (7/30-day lookbacks)

    - ■ `churn_enriched` with segment/market/onboarding, tenure, activity, and `days_since_last_use`

  - ○ Creates views for top reasons, monthly churn, and usage-by-segment

  - ○ **Exports**: `kpi_churn_trend.csv`, `top_reasons.csv`, `reasons_by_segment.csv`, `risk_indicators.csv`

- ● `notebooks/otter_churn_starter.py`

  - ○ Runs the SQL end-to-end and prints out the produced files

  - ○ Shows a sample of model-ready features (you can extend to a full logistic regression if your dataset has both churned and retained labels)

- ● `requirements.txt` with `duckdb`, `pandas`, `pyarrow`, `polars`, `matplotlib`, `scikit-learn`
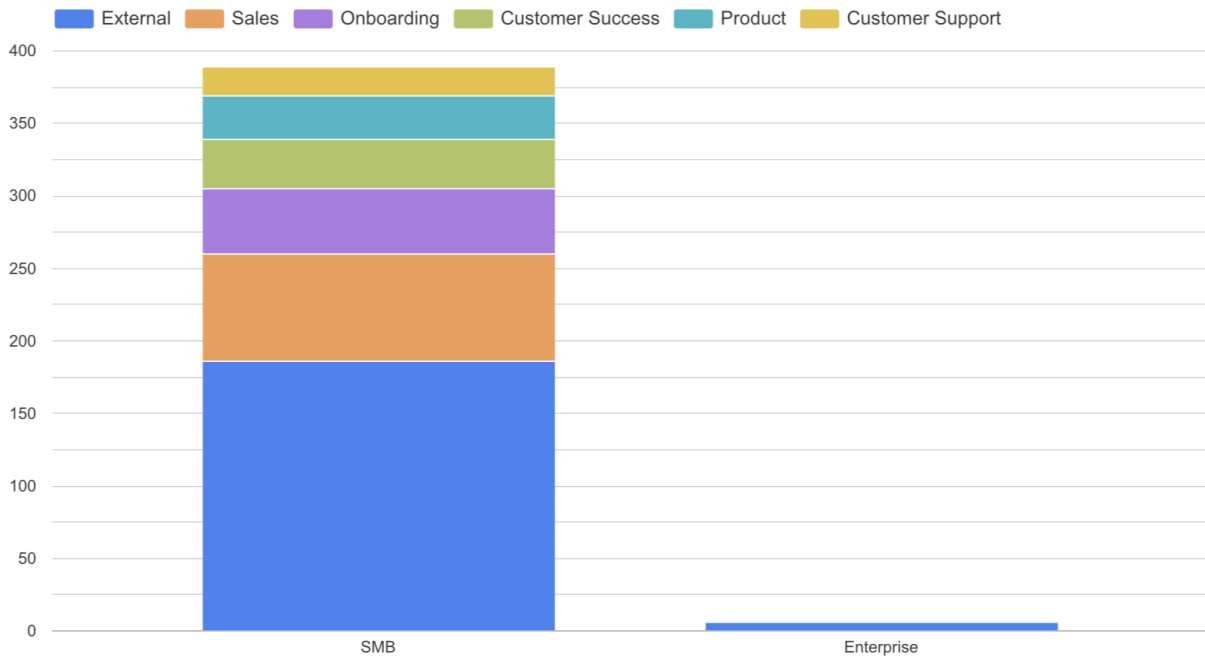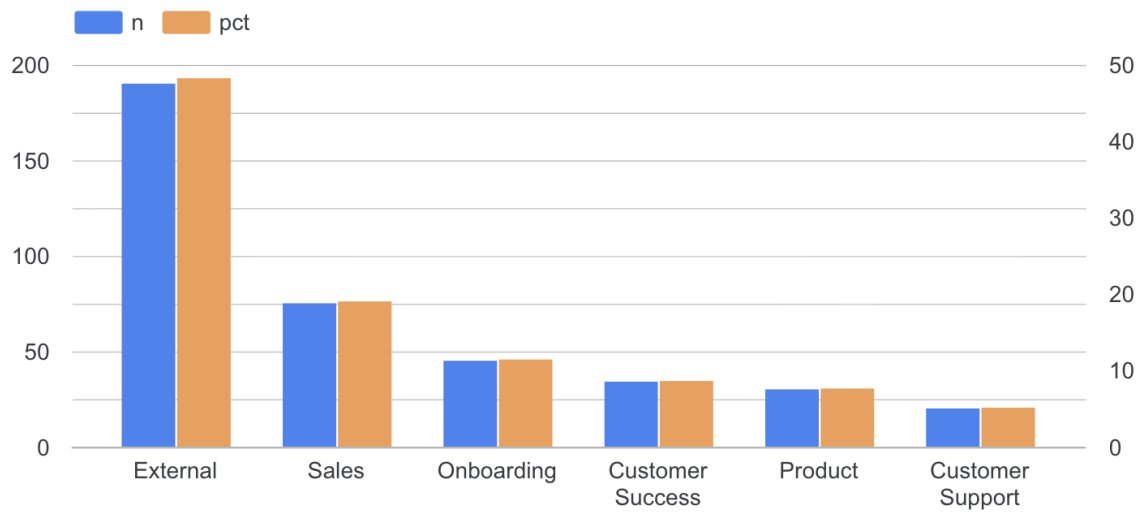
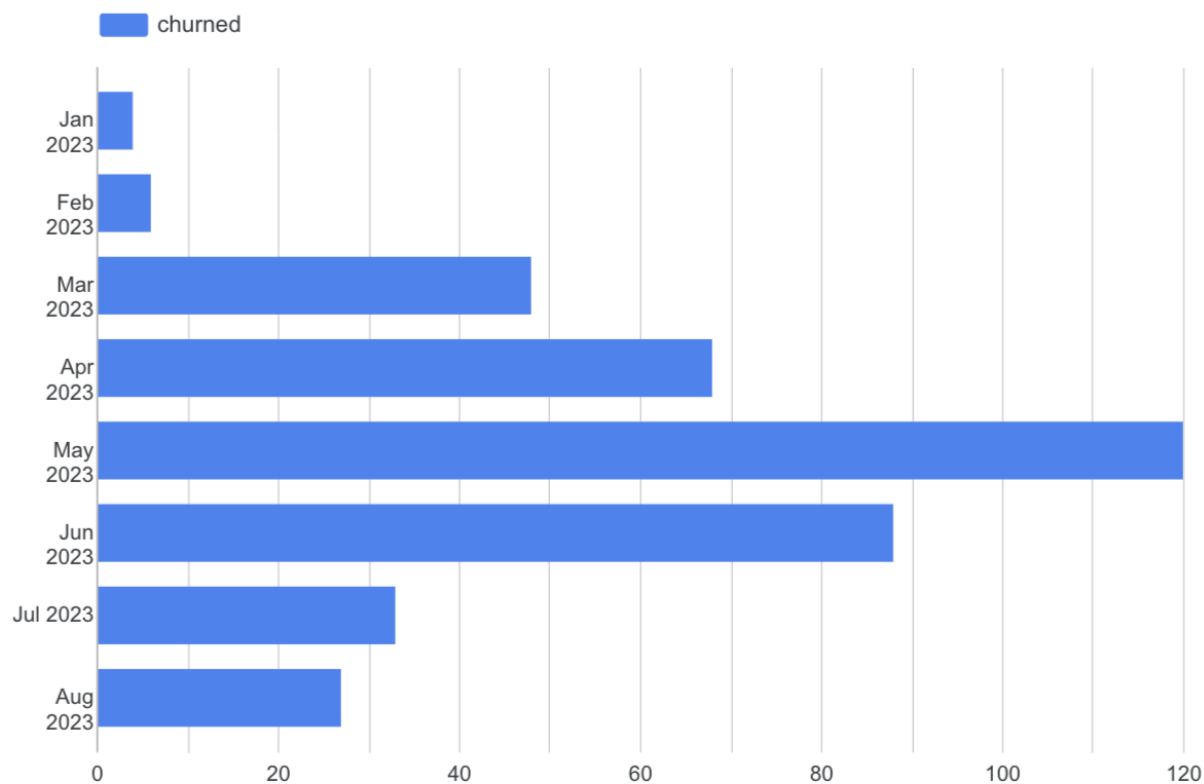- ● `README.md` with the condensed instructions

Visualizations -

Sentiment Distribution (Weighted Text)

Average sentiment: −0.000505568390759624

Top Keywords (Weighted: sec×3 + pri×2 + notes×1)

Word Cloud — Weighted Reasons + Notes

Top Primary Reasons

churn_primary_reason	
Customer Support	
Customer Success	
Product	
Onboarding	
Sales	
External	

(Horizontal axis: 0, 50, 100, 150, 200, 250)

Top Secondary Reasons

churn_secondary_reason	
Switching to a competitor	
Customer unwilling to provide feedback	
Other customer success issue	
Incorrect expectations	
Other product-related issues	
Setup not complete prior to first payment	
Dormant if pre-active	
Other Sales-related issues	
Poor support experience	
Poor product fit	
Resto closing	
Other onboarding-related issues	
Low Volume of Sales in Delivery	
Churn Automation - DO NOT USE	
Other external issues	

(Horizontal axis: 0, 20, 40, 60, 80, 100, 120)

churned

Jan 2023	
Feb 2023	
Mar 2023	
Apr 2023	
May 2023	
Jun 2023	
Jul 2023	
Aug 2023	

0   20   40   60   80   100   120

Average Days Since Last Used
1.26

Average Events (30d)
25.18

Average Active Users (30d)
21.27